

HANOI UNIVERSITY OF  
SCIENCE AND TECHNOLOGY

# A Scalable Lakehouse for Financial Analytics

Big Data Storage and Processing Capstone Project

---

Trinh Hoang Anh  
Bui Khac Chien  
Bach Nhat Minh

Tran Quang Minh  
Vo Ta Quang Nhat

Supervised by: Dr. Tran Viet Trung

# Table of Contents

1. Introduction & Problem Statement
2. The Architectural Blueprint
3. Infrastructure & Deployment
4. Notable Lessons Learned
5. Conclusion & Key Achievements

# 1. The Challenge: Taming Financial Data

Financial markets generate data at extreme volume, velocity, and variety, creating significant engineering challenges that demand a unified architecture.

## Heterogeneous Sources

Integrating real-time streams (Finnhub quotes) with massive historical datasets (Stooq prices, GDELT global news).

## Dual Analytical Demands


Supporting low-latency dashboards for real-time monitoring **and** deep, historical analysis for ML models.

## Data Integrity & Governance


Ensuring consistency, quality, and traceability across a massive, evolving dataset from raw ingestion to aggregated insights.

**Project Mission:** To design and implement a scalable Big Data Lakehouse capable of ingesting, processing, and analyzing financial data in both real-time and batch modes.

## 2. Unified Lambda Architecture



## 2a. Ingestion: A Hybrid, Purpose-Built Strategy



### Batch Layer (High Throughput)

Sources: GDELT, Stooq.



Principle: A 'Direct-to-Lake' approach bypasses relational databases to maximize throughput and avoid I/O bottlenecks for bulk historical data.

### Speed Layer (Low Latency & Integrity)

Source: Finnhub API.



Principle: Uses PostgreSQL as a transactional buffer for data persistence deduplication. Debezium's Change Data Capture on the database logs creates a real-time, event-driven stream without impacting the source.

## 2b. Processing: The Medallion Lakehouse on Spark & Iceberg

### **Bronze (Raw)**

Immutable, raw data from sources (GDELT, Stooq) stored as-is. Partitioned by date for efficient writes. The single source of truth for reprocessing.

### **Silver (Cleaned & Enriched)**

Contains validated, structured, and deduplicated data. Key transformations include schema enforcement, temporal normalization (UTC to EST), and linking news events to specific stock tickers.

### **Gold (Aggregated)**


Business-level aggregates optimized for analytics and ML. Examples: daily sentiment scores, news volume, and volatility metrics per ticker.

### **Key Enabler - Apache Iceberg**

Stored on MinIO, Iceberg provides ACID transactions, schema evolution, and time-travel directly on the data lake, combining data warehouse reliability with data lake scalability.

## 2c. Serving: Unifying Batch and Stream with ClickHouse

ClickHouse provides a single, high-performance OLAP access point for both deep historical data and low-latency real-time streams.



### Integration Strategies:

- **Batch Data:** Utilizes the Iceberg Table Engine for **zero-copy queries** directly on Parquet files in MinIO. Aggregated data is also loaded into 'MergeTree' tables for maximum dashboard performance.
- **Stream Data:** Ingests high-throughput streams from Kafka into optimized 'MergeTree' tables designed for time-series queries.

**Outcome:** Sub-second query performance for interactive dashboards and immediate visibility of market data.

### 3. The Stack: A Containerized & Cloud-Native Foundation

The entire system is built on a containerized, cloud-agnostic infrastructure, enabling reproducible and scalable deployments.


Deployment & Orchestration



Processing & Storage



Machine Learning



Workflow & Metadata



Streaming & Database



Core Principles: Infrastructure as Code (IaC) • Scalability • Fault Tolerance

## 4. Notable Lesson #1: Hybrid Ingestion is a Necessity, Not a Choice




### Problem

Clash Grotesk Medium

A “one-size-fits-all” ingestion approach created severe bottlenecks.

Routing massive batch files (GDEL) through PostgreSQL caused I/O overhead, while writing API streams directly to the lake risked data duplication.



#### Direct-to-Lake

Spark writes directly to MinIO, bypassing transactional overhead for maximum throughput.

### Solution

Clash Grotesk Medium



#### Transactional Buffering

API -> Postgres -> CDC, using the database for its strengths in enforcing integrity and deduplication before streaming.

### Key Takeaway

Don't put a database in the middle unless you have a transactional reason. Use raw object storage for bulk throughput.


## 4. Notable Lesson #2: Separate Processing from Serving for Low-Latency BI

### Problem



Dashboards were sluggish and timed out. Connecting Superset directly to Spark was slow (high latency), and using PostgreSQL for large-scale aggregations was inefficient (row-oriented).

### Solution



#### Dedicated OLAP Serving Layer


Implemented **ClickHouse** as a dedicated **OLAP serving layer**. Its columnar storage and vectorized query execution are purpose-built for the fast aggregation queries required by BI dashboards, delivering sub-second responses.

### Key Takeaway

Spark is for throughput (ETL processing). ClickHouse is for latency (interactive serving). Use the right engine for the query pattern. Satoshi Regular

## 4. Notable Lesson #3: Data Schemas as Contracts Enable Parallel Work

### Problem




The visualization team was blocked, waiting for the data engineering pipeline to be completed. This forced a slow, sequential “waterfall” process.

### Solution

#### Traditional Waterfall



#### Contract-First Development



Adopted **Contract-First Development**. The teams agreed on the final “Gold” layer data schema (column names, types, format) upfront. The visualization team built dashboards using mock data matching this contract, while the DE team built the pipeline to produce it. Integration required zero changes to the dashboard logic.

### Key Takeaway

A well-defined data contract decouples teams, is the foundation for agile development in data projects, and drastically reduces integration risk.

## 4. Notable Lesson #4: Robust Deployment Requires Infrastructure as Code

### Problem




Managing 15+ Kubernetes resources with manual `kubectl apply` commands and shell scripts was brittle, error-prone, and lacked versioning.

### Key Takeaway

Helm treats infrastructure as versioned code (IaC), preventing configuration drift and making complex deployments manageable, resilient, and repeatable.

### Solution




Migrated all Kubernetes YAML manifests into a unified **Helm chart**. This templated the entire application, centralizing configuration in a single `values.yaml` file and enabling versioned, atomic deployments and rollbacks with a single command.

## 5. Conclusion: An End-to-End Financial Data Platform

We successfully designed and implemented a scalable, end-to-end Big Data Lakehouse that addresses the core challenges of complex financial analysis.

- ✓ **Unified Architecture:** Integrated batch and stream processing via a Lambda pattern.
- ✓ **Data Quality & Reliability:** Ensured curated, reliable data at data with the Medallion model on Apache Iceberg.
- ✓ **High-Performance Analytics:** Delivered low-latency dashboards through a dedicated OLAP serving layer (ClickHouse).
- ✓ **Robust & Reproducible System:** Built on a modern, containerized infrastructure managed with Infrastructure as Code (Helm).



Simplified Architecture Overview